

## Example: Bellman Equation for Q-Function

Consider a simple grid world with three states (S1, S2, and the goal state G) and two actions (left and right). The agent receives a reward of  $-1$  for each step and a reward of  $+10$  for reaching the goal state. The discount factor  $\gamma$  is set to  $0.9$ .

The Bellman equation for the Q-function  $Q(s, a)$  of a state  $s$  and action  $a$  is given by:

$$Q(s, a) = R(s, a) + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q(s', a')$$

where:

- $R(s, a)$  is the immediate reward received after taking action  $a$  in state  $s$ ,
- $\gamma$  is the discount factor,
- $P(s'|s, a)$  is the probability of transitioning to state  $s'$  from state  $s$  after taking action  $a$ .

Let's calculate the Q-values for each state-action pair:

1. **Initialization:** Start with initial Q-values of zero for all state-action pairs.

State	Action (left)	Action (right)
S1	0	0
S2	0	0
G	0	0

2. **Update Q-Value for State S1 and Action Left:**

$$R(S1, \text{left}) = -1 \quad (\text{immediate reward for moving left})$$

$$\max_{a'} Q(S2, a') = \max(0, 0) = 0 \quad (\text{maximum Q-value for the next state S2})$$

$$Q(S1, \text{left}) = -1 + 0.9 \times 0 = -1$$

3. **Update Q-Value for State S1 and Action Right:**

$$R(S1, \text{right}) = -1 \quad (\text{immediate reward for moving right})$$

$$\max_{a'} Q(G, a') = \max(0, 0) = 0 \quad (\text{maximum Q-value for the next state G})$$

$$Q(S1, \text{right}) = -1 + 0.9 \times 0 = -1$$

4. **Update Q-Value for State S2 and Action Left:**

$$R(S2, \text{left}) = -1 \quad (\text{immediate reward for moving left})$$

$$\max_{a'} Q(G, a') = \max(0, 0) = 0 \quad (\text{maximum Q-value for the next state G})$$

$$Q(S2, \text{left}) = -1 + 0.9 \times 0 = -1$$

5. **Update Q-Value for State S2 and Action Right:**

$$R(S2, \text{right}) = -1 \quad (\text{immediate reward for moving right})$$

$$\max_{a'} Q(G, a') = \max(0, 0) = 0 \quad (\text{maximum Q-value for the next state G})$$

$$Q(S2, \text{right}) = -1 + 0.9 \times 0 = -1$$

6. **Update Q-Value for State G and Actions:**

$$R(G, \text{left}) = 10 \quad (\text{reward for reaching the goal state})$$

$$R(G, \text{right}) = 10 \quad (\text{reward for reaching the goal state})$$

$$Q(G, \text{left}) = 10 + 0.9 \times 0 = 10$$

$$Q(G, \text{right}) = 10 + 0.9 \times 0 = 10$$

After updating all state-action pairs, the Q-values are as follows:

State	Action (left)	Action (right)
<i>S1</i>	-1	-1
<i>S2</i>	-1	-1
<i>G</i>	10	10

The final Q-values represent the expected cumulative rewards the agent can achieve from each state-action pair following an optimal policy.