

Solving the Frozen Lake Problem with Value Iteration

Here's the Python code to solve the Frozen Lake problem using the value iteration algorithm:

```
1 import gym
2 import numpy as np
3
4 # Create the Frozen Lake environment
5 env = gym.make('FrozenLake-v1')
6
7 # Initialize the value function V(s) to zeros
8 V = np.zeros(env.observation_space.n)
9
10 # Set the discount factor
11 gamma = 0.9
12
13 # Set the convergence threshold
14 epsilon = 1e-6
15
16 # Value iteration algorithm
17 while True:
18     delta = 0
19     for s in range(env.observation_space.n):
20         v = V[s]
21         # Calculate the new value for state s using the
22         # Bellman equation
23         V[s] = np.max([sum([p*(r + gamma*V[s_]) for p, s_, r
24         , _ in env.P[s][a]]) for a in range(env.action_space.n)])
25         delta = max(delta, np.abs(v - V[s]))
26     if delta < epsilon:
27         break
28
29 # Extract the optimal policy from the optimal value function
30 policy = np.array([np.argmax([sum([p*(r + gamma*V[s_]) for p
31 , s_, r, _ in env.P[s][a]]) for a in range(env.
32 action_space.n)]) for s in range(env.observation_space.n)
33 ])
34
35 # Display the optimal policy
36 print("Optimal policy:")
37 print(policy.reshape((4, 4)))
38
39 # Test the learned policy
40 total_reward = 0
41 num_episodes = 100
42 for _ in range(num_episodes):
43     state = env.reset()
44     done = False
```

```

40     while not done:
41         action = policy[state]
42         state, reward, done, _ = env.step(action)
43         total_reward += reward
44
45     # Calculate the average reward
46     average_reward = total_reward / num_episodes
47     print(f"Average reward over {num_episodes} episodes: {
        average_reward}")

```

This Python code uses the OpenAI Gym library to create the Frozen Lake environment and implements the value iteration algorithm to find the optimal policy. The algorithm iteratively updates the value function until convergence, then extracts the optimal policy from the optimal value function and tests it over multiple episodes to calculate the average reward.